

電腦視覺與深度學習

(Computer Vision and Deep Learning)

Homework 2

TA: 胡滋紘

Gmail: nckubot65904@gmail.com

Office Hour: 14:00~16:00, Mon.

10:00~12:00, Fri.

At CSIE 9F Robotics Lab.

Notice (1/2)

- Copying homework is strictly prohibited!! **Penalty: Both individuals will receive a score of 0!!**
- Due date => **09:00:00, 2023/12/21 (Thu.)**
 - Do not submit late**, or the following points will be deducted:
 - Submit within seven days after the deadline, and your score will be reduced by half.
 - If you submit after this period, you will receive a score of 0.
- You must **attend the demonstration**, otherwise your score will be 0. The demonstration schedule **will be announced on NCKU Moodle**.
- You must **create GUI**, otherwise your point will be **deducted**.
- Upload to => **140.116.154.28 -> Upload/Homework/Hw2**
 - **User ID: cvdl2023 Password: RL2023cvdl**
- Format
 - Filename: **Hw2_StudentID_Name_Version.rar**
 - **Ex: Hw2_F71234567_林小明_V1.rar**
 - If you want to update your file, you should update your version to be V2,
 - **Ex: Hw2_F71234567_林小明_V2.rar**
 - Content: **Project folder** *(Excluding the pictures)
 - *Note: Remove your “Debug” folder to reduce file size.

Notice (2/2)

- Python (recommended):
 - Python 3.8 (<https://www.python.org/downloads/>)
 - **Opencv-contrib-python (3.4.2.17)**
 - Matplotlib 3.7.3
 - UI framework: pyqt5 (5.15.10)
 - Pytorch 2.1.0
 - Torchvision 0.16.0
 - Torchsummary 1.5.1
 - Tensorboard 2.14.0
 - Pillow 10.1.0

Assignment scoring (Total: 100%)

1. (20%) Background Subtraction (出題 : Chen)

2. (20%) Optical Flow (出題 : Jimmy)

2.1 (10%) Preprocessing

2.2 (10%) Video tracking

3. (20%) PCA - Dimension Reduction (出題 : Zhong)

4. (20%) Training a MNIST Classifier Using VGG19 with BN (出題 : Shang)

4.1 (6%) Load Model and Show Model Structure.

4.2 (6%) Show Training/Validating Accuracy and Loss.

4.3 (8%) Use the Model with Highest Validation Accuracy to Run Inference, Show the Predicted Distribution and Class Label.

5. (20%) Train a Cat-Dog Classifier Using ResNet50

(出題 : Shan)

5.1 (4%) Load the dataset and resize images

5.2 (4%) Plot class distribution of training dataset

5.3 (4%) Show the structure of ResNet50 model

5.4 (4%) Improve ResNet50 with Random-Erasing and Compare the accuracies of 2 ResNet50 models on validation dataset

5.5 (4%) Use the trained model to run inference and show the predicted class label

* Don't fix your image and video path
(There is another dataset for demonstration)

Load image and video please use the following function to read the path.

[QFileDialog.getOpenFileName](#)



Assignment scoring (Total: 100%)

- Use one UI to present 5 questions.

Load Image

Load Video

1. Background Subtraction

1. Background Subtraction

2. Optical Flow

2.1 Preprocessing

2.2 Video tracking

3. PCA

3. Dimension Reduction

4. MNIST Classifier Using VGG19

1. Show Model Structure

2. Show Accuracy an Loss

3. Predict

4. Reset

5. ResNet50

Load Image

5.1. Show Images


5.2. Show Model Structure

5.3. Show Comprasion

5.4. Inference


1. Background Subtraction (20%)

- Given: a traffic video : “traffic.mp4”
- Q: Please remove background and show the result (Result should be like the Demo below)
 - 1) Load Video from File Dialog
 - 2) Create subtractor using `cv2.createBackgroundSubtractorKNN(history,dist2Threshold,detectShadows=True)`
 For each frame in video:
 - 3.1) Blur frame using `cv2.GaussianBlur(frame, (5, 5), 0)`
 - 3.2) Get background mask (**M**) by `subtractor.apply`
 - 3.3) Generate Frame (**R**) with only moving object by `cv2.bitwise_and`



rgb frame

3.1) ~
3.2)

Foreground mask
(**M**)

3.3)

Result (**R**)

Demo (show result when frame number >4)

(**M**)

(**R**)



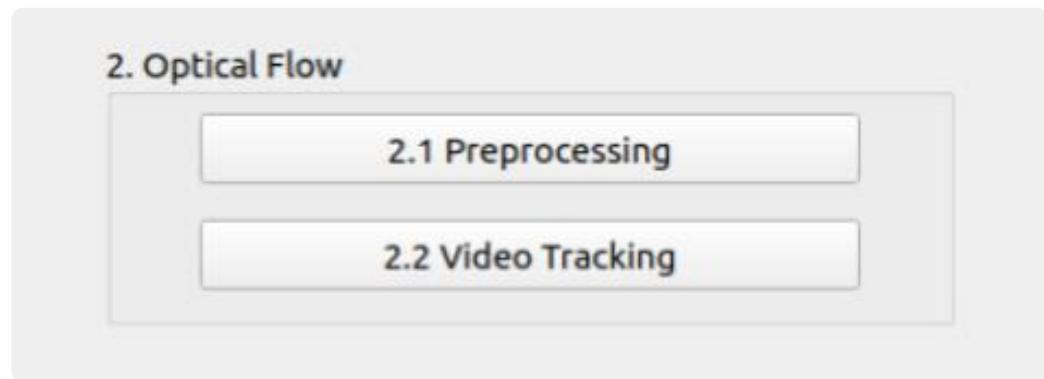
2. (20%) Optical Flow

(出題：Jimmy)

2.1 (10%) Preprocessing

2.2 (10%) Video tracking

UI demo:




2.1 Preprocessing (10%)



(出題 : Jimmy)

❑ Given a video: “optical_flow.mp4”

❑ Q: Click the button 2.1 to detect the point at the bottom of the doll's nose in **1st Frame** using `cv2.goodFeaturesToTrack` and show the point with a **red cross mark** using `cv2.line`



- Parameters in `goodFeaturesToTrack` should be adjusted:
 1. `maxCorners`: 1
 2. `qualityLevel` : 0.3
(quality = $\min(\lambda_1, \lambda_2)$, λ : eigenvalue from the Hessian matrix)
 3. `minDistance`: 7
 4. `blockSize`: 7



2.2 Video tracking (10%)

(出題 : Jimmy)

❑ Q: Click button 2.2 to:

- 1) (5%) Track the point (detected in 2.1) on the whole video using OpenCV function [cv2.calcOpticalFlowPyrLK](#).
- 2) (5%) Display the [trajectory](#) of the tracking point throughout the video using [cv2.line](#). Pick a **highly visible color**. Ex: Yellow (0,100,255)

Demo videos:



❑ Tool site:

1. [Load video](#)
2. [Optical flow](#)
3. [Optical flow Tutorial](#)

3. (20%) PCA – Dimension Reduction

- ❑ Given: A RGB image “logo.jpg”
- ❑ Q: Using **PCA (Principal components analysis)** to do dimension reduction on given image, find **the minimum components** that **reconstruction error less or equal to 3.0**
 - 1) Convert RGB image to gray scale image, image shape will be (w,h).
 - 2) Normalize gray scale image from [0,255] to [0,1]
 - 3) Use PCA to do dimension reduction **from min(w,h) to n**, then reconstruct the image.
 - 4) Use MSE to compute reconstruction error, and find minimum **n** that error value less or equal to 3.0. **Print out the n value.** (10%)
 - 5) Plot the **gray scale image** and the **reconstruction image** with **n** components. (10%)
- ❑ Hint: Use PCA from python library: sklearn.decomposition

Mean Square Error (MSE) Formula:

$$\frac{1}{n_{pixels}} * \sum_{i=1}^{n_{pixels}} (\vec{x}_i - \vec{y}_i)^2, \text{ where}$$

\vec{x}_i is the gray value of pixel in original image (gray scale image)

\vec{y}_i is the gray value of pixel in reconstruction image



logo.jpg



Gray scale image



Reconstruction image (i.e. n=50)

4. Training a MNIST Classifier Using VGG19 with BN (20%)

(出題：Shang)

4.1 Load Model and Show Model Structure. (6%)

4.2 Show Training/Validating Accuracy and Loss. (6%)

4.3 Use the Model with Highest Validation Accuracy to Run Inference, Show the Predicted Distribution and Class Label. (8%)





Figure: GUI example

4.0 Training a MNIST Classifier Using VGG19 with BN (20%)

(出題 : Shang)

■ Requirement

- 1) Train VGG19 model with batch normalization (BN) using **PyTorch**.
- 2) Download dataset using torchvision.datasets.MNIST() ([tutorial](#))
 - Training data: 60000 images
 - Validation data: 10000 images
 - Resize image to (32, 32)
- 3) Parameters
 - At least **30 epochs**.
 - **Cross entropy loss**
 - **Adam** optimizer
- 4) Record training/validation loss and accuracy in **.jpg** or **.png** format.
- 5) In the submitted file, you need to include
 - Weight file for VGG19 with BN in **.pth** format. (File size is approximately 540MB)
 - Figure of training/validating loss and accuracy in **.jpg** or **.png** format.
 - Code for your GUI program
 - Code for model training.
- 6) **Please do not include image data in the submitted file.**



R. Reference

- 1) [VGG19](#)
- 2) [Batch Normalization](#)

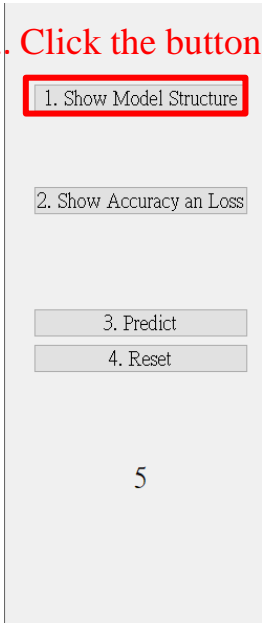
4.1 Show the Structure of VGG19 with BN (6%)

1. Click the button “1. Show Model Structure”
2. Show the VGG19 with BN model on terminal using `torchsummary.summary()`.

(出題 : Shang)

The -1 indicates that the actual size of batch size can vary.

1. Click the button.



Layer (type)	Feature map shape (Batch, Channels, Height, Width)	Num. of param.
BatchNorm2d-38	[-1, 512, 4, 4]	1,024
ReLU-39	[-1, 512, 4, 4]	0
MaxPool2d-40	[-1, 512, 2, 2]	0
Conv2d-41	[-1, 512, 2, 2]	2,359,808
BatchNorm2d-42	[-1, 512, 2, 2]	1,024
ReLU-43	[-1, 512, 2, 2]	0
Conv2d-44	[-1, 512, 2, 2]	2,359,808
BatchNorm2d-45	[-1, 512, 2, 2]	1,024
ReLU-46	[-1, 512, 2, 2]	0
Conv2d-47	[-1, 512, 2, 2]	2,359,808
BatchNorm2d-48	[-1, 512, 2, 2]	1,024
ReLU-49	[-1, 512, 2, 2]	0
Conv2d-50	[-1, 512, 2, 2]	2,359,808
BatchNorm2d-51	[-1, 512, 2, 2]	1,024
ReLU-52	[-1, 512, 2, 2]	0
MaxPool2d-53	[-1, 512, 1, 1]	0
AdaptiveAvgPool2d-54	[-1, 512, 7, 7]	0
Linear-55	[-1, 4096]	102,764,544
ReLU-56	[-1, 4096]	0
Dropout-57	[-1, 4096]	0
Linear-58	[-1, 4096]	16,781,312
ReLU-59	[-1, 4096]	0
Dropout-60	[-1, 4096]	0
Linear-61	[-1, 10]	40,970

Total params: 139,622,218
Trainable params: 139,622,218
Non-trainable params: 0

Input size (MB): 0.01
Forward/backward pass size (MB): 7.55
Params size (MB): 532.62
Estimated Total Size (MB): 540.18

Figure: the Structure of VGG19 with BN

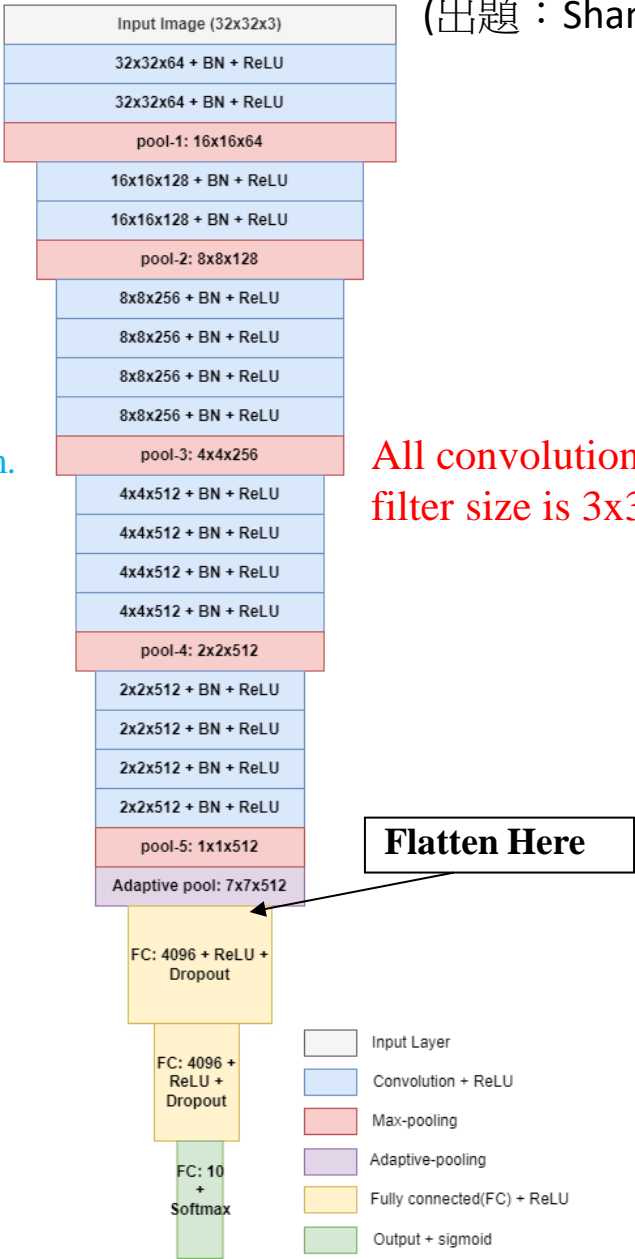


Figure: VGG19 with BN model structure

4.2 Show Training/Validating Accuracy and Loss (6%)


1. At home:

(出題 : Shang)

- 1) Download the training and validation datasets. ([tutorial](#))
- 2) Training and validating VGG19 with BN **at least 30 epochs** at home ([tutorial](#)) and record the training/validating accuracy and loss in each epoch ([tutorial](#)).
- 3) If your validation accuracy is low, you can try
 - Adjust the **learning rate** of the optimizer.
 - Change the **data augmentation** techniques used.
- 4) Save weight file with **highest** validation accuracy .
- 5) Use [matplotlib.pyplot.plot\(\)](#) to create a line chart for the **training and validating loss and accuracy** values and save the figure.

2. When the demo:

- (1) Click the button “2. Show Accuracy and Loss”
- (2) Show the **saved figure** of Training/Validating loss and accuracy in a new window




4.3 Use the Model with Highest Validation Accuracy to Run Inference, Show the Predicted Distribution and Class Label. (8%) (出題 : Shang)

1. Load the model with **highest validation accuracy** which trained at home.
2. Draw a number on graffiti board using **mouse**. ([tutorial](#))
 - Background: **black**
 - Pen: **white**
3. Click the button “**3. Predict**” to run inference on the image you drew.
 - Show the predicted class label on the GUI.
 - Show the probability distribution of model predictions using a histogram in a new window.
4. Click the button “**4. reset**” to clear the graffiti board.




3. Probability of each class



4.3 Use the Model with Highest Validation Accuracy to Run Inference, Show the Predicted Distribution and Class Label. (6%) (出題 : Shang)

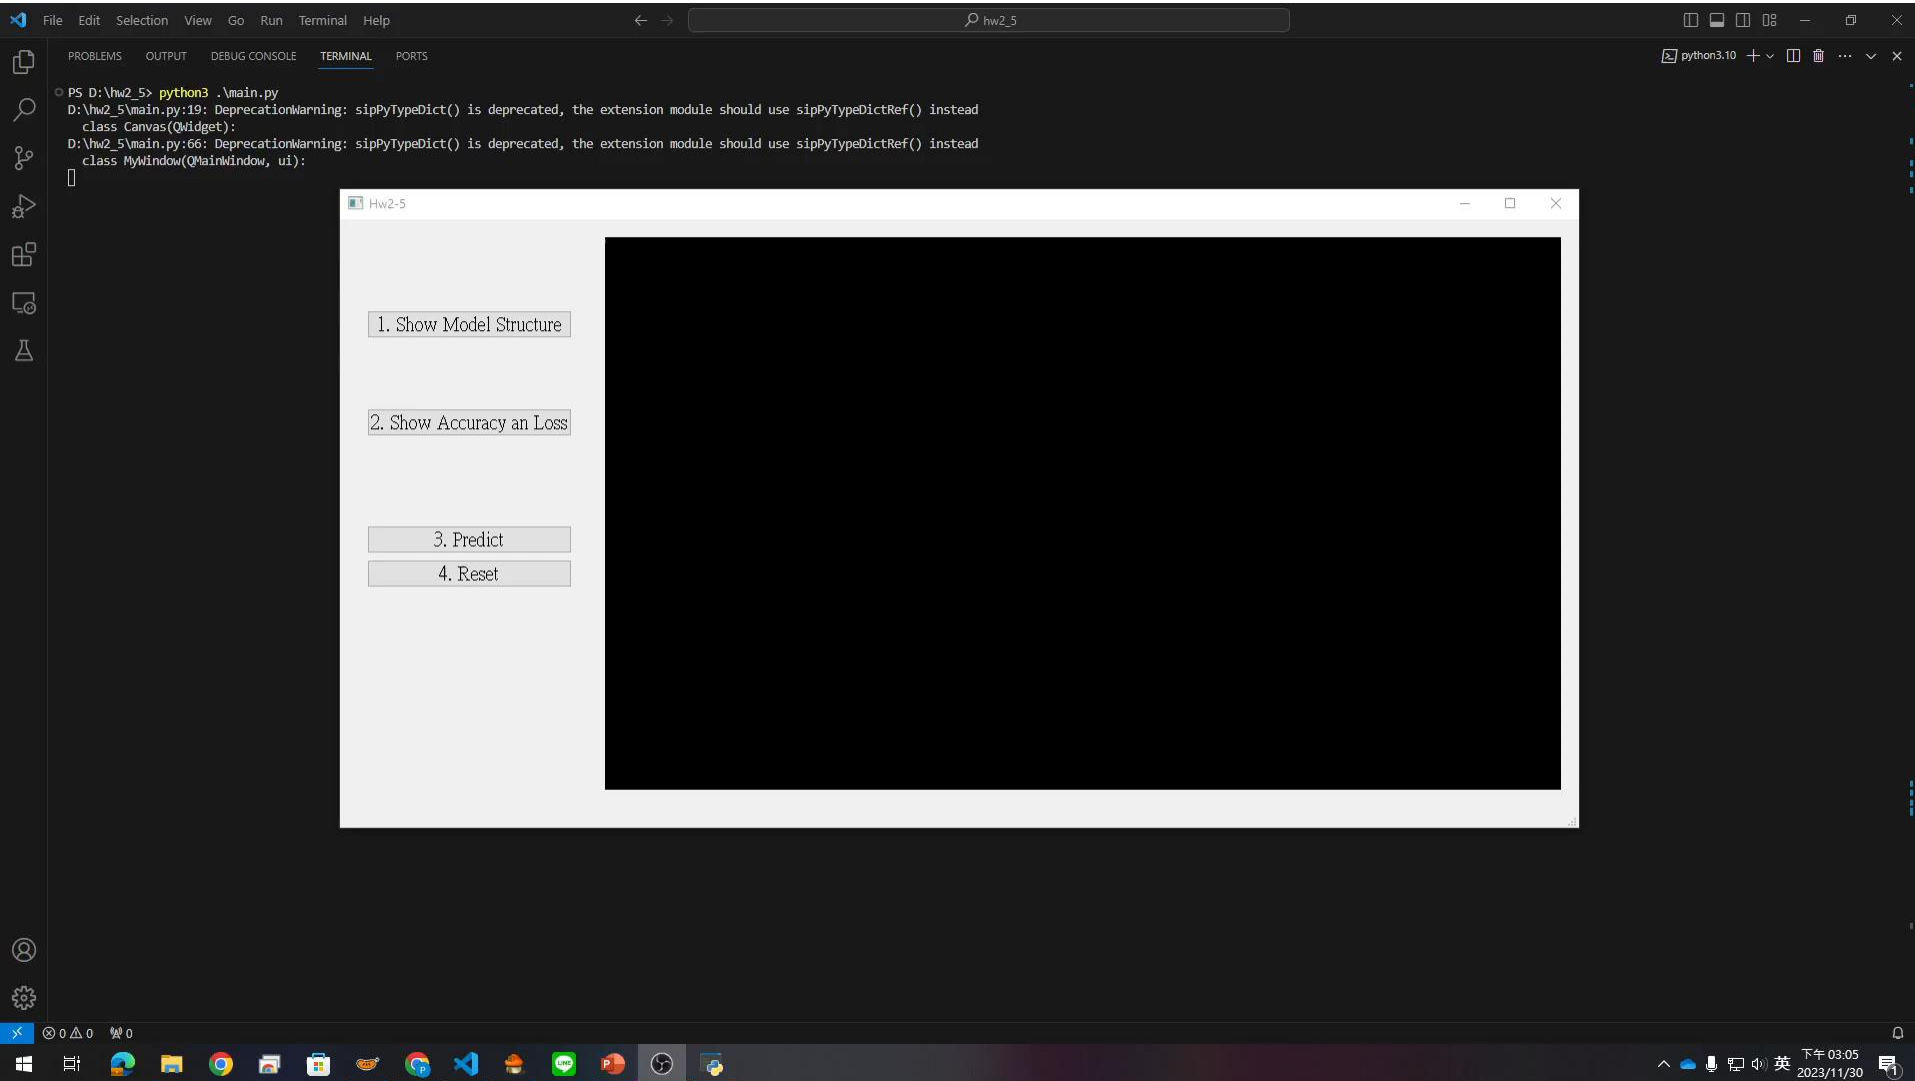
- The probability distribution of model prediction using a histogram.



4. Training a MNIST Classifier Using VGG19 – Example Video

(出題：Shang)

- This is an example illustrating the objectives from 4.1 ~ 4.3.



5. Train a Cat-Dog Classifier Using ResNet50 (20%) (出題：Shan)

5.1 (5%) Load the dataset and resize images

5.2 (5%) Show the structure of ResNet50 model

5.3 (5%) Improve ResNet50 with Random-Erasing and Compare the accuracies of 2 ResNet50 models on validation dataset

5.4 (5%) Use the trained model to run inference and show the predicted class label

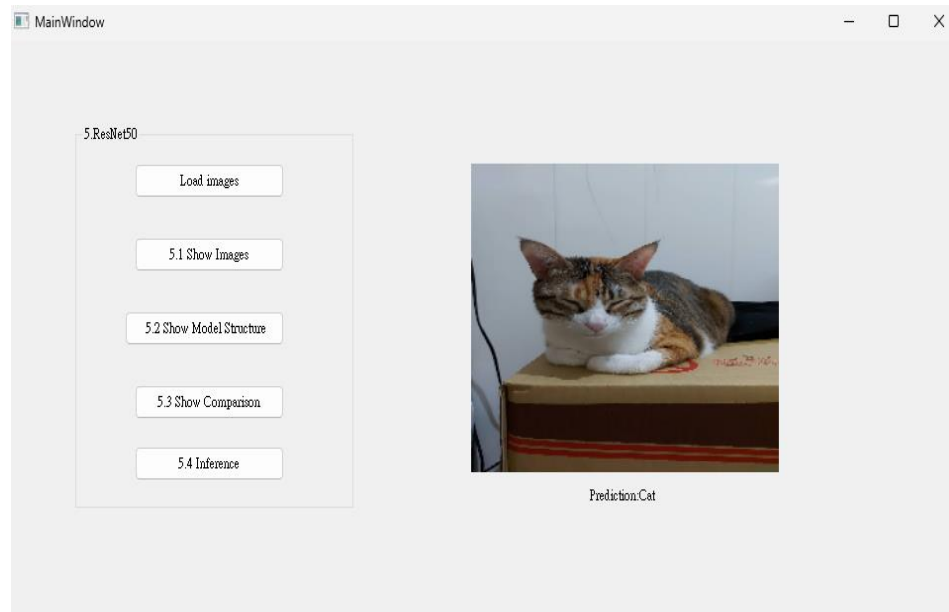


Figure: GUI Example

5.0 Train a Cat-Dog Classifier Using ResNet50

(出題：Shan)

1. Objective

- 1) Learn how to train a ResNet50 model to **classify images** of cats and dogs using PyTorch ([tutorial](#))

2. Download Cats and Dogs Dataset from FTP

- 1) Data type: JPG images
- 2) 2 classes: Cat and Dog
- 3) Datasets
 - (1) Training dataset: 16,200 JPG images in total.
 - (2) Validation dataset: 1,800 JPG images in total.
 - (3) Inference dataset: 10 JPG images in total.

It is for **testing the inference function** in your GUI program.

3. In the submitted file

- 1) Organize the files in this structure:

```
Hw2_StudentID_Name_Version // project folder
|-- model                  // folder to put trained models
|-- inference_dataset
    |-- Cat
    |-- Dog
|-- main.py                // codes for your GUI program
|-- train.py               // codes for model training
|-- ...                    // other files or folders you need
```

Notice: Please include the inference dataset in your homework file.

R. Reference

- 1) [Deep Residual Learning for Image Recognition](#)
- 2) [Kaggle Cats and Dogs Dataset](#)

5.1 (5%) Load the dataset and resize images

(出題：Shan)

1) At home:

(1) Load the **inference dataset**

→ Hint:

(a) PyTorch ([tutorial](#)): `torch.utils.data.Dataset`

(2) Resize images to **224×224×3c** (RGB)

→ Hint:

(a) PyTorch ([tutorial](#)): `torchvision.transform`

(3) Click the button **“1. Show Images”**

(4) Get 1 image from **each class** in the inference dataset

(5) Show images in a **new window**

→ Hint: use `matplotlib.pyplot` functions to show images ([tutorial](#)):

(a) `figure()`

(b) `imshow()`

(c) `subplot()`

(d) `title()`

2) When the demo:

(1) Click the button **“1. Show Images”**

(2) Show images in a **new window**

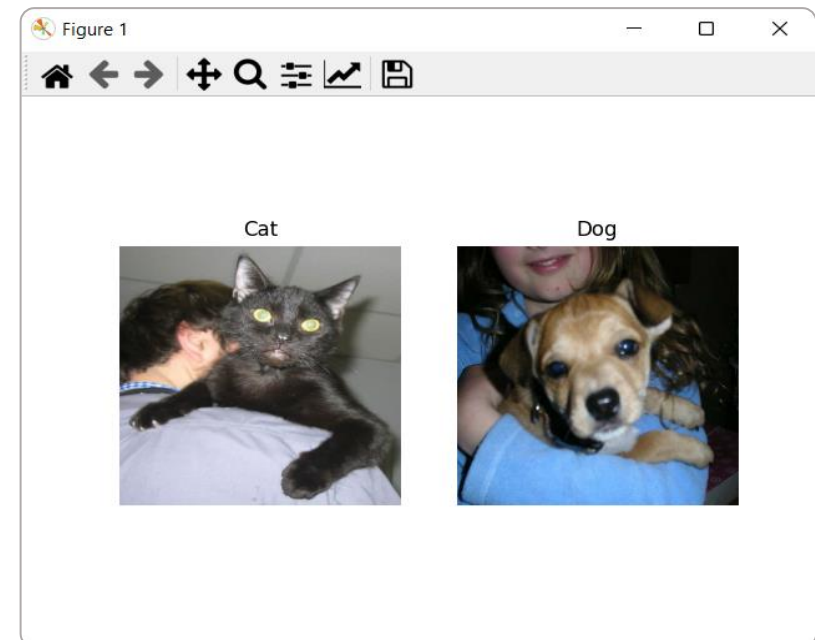
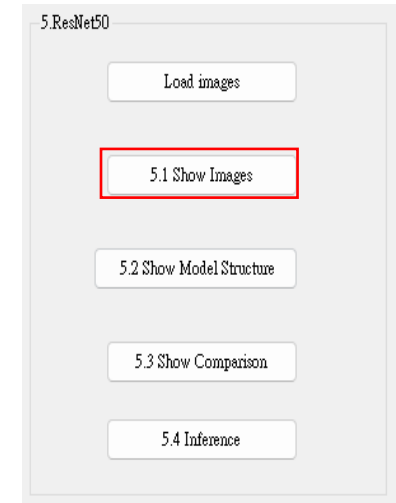


Figure: 1 image from each class
Notice: this is an example, the images might differ

5.2 (5%) Show the structure of ResNet50 model

(出題：Shan)

1) At home:

(1) Build a ResNet50 model

→ Hint:

(a) PyTorch: `torchvision.models.resnet50()`

(2) Replace the output layer to a FC (Fully Connected) layer of **1 node** with a **Sigmoid** activation function

→ Hint:

(a) PyTorch ([tutorial](#)): `torch.nn.Linear(2048, 1), torch.nn.Sigmoid`

If the class label of Cat is 1, the output value (range: 0 ~ 1) should be **close to 1** for cat images, and vice versa.

(3) Run the function to show the structure **in the terminal**

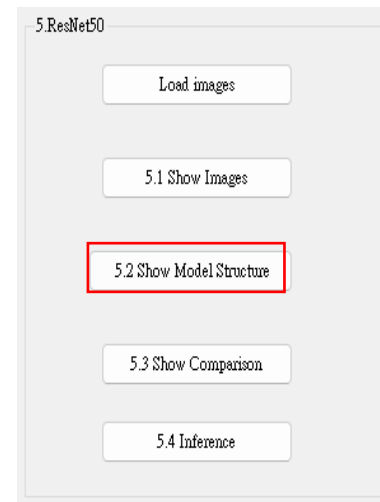
→ Hint:

(a) PyTorch: `torchsummary`

2) When the demo:

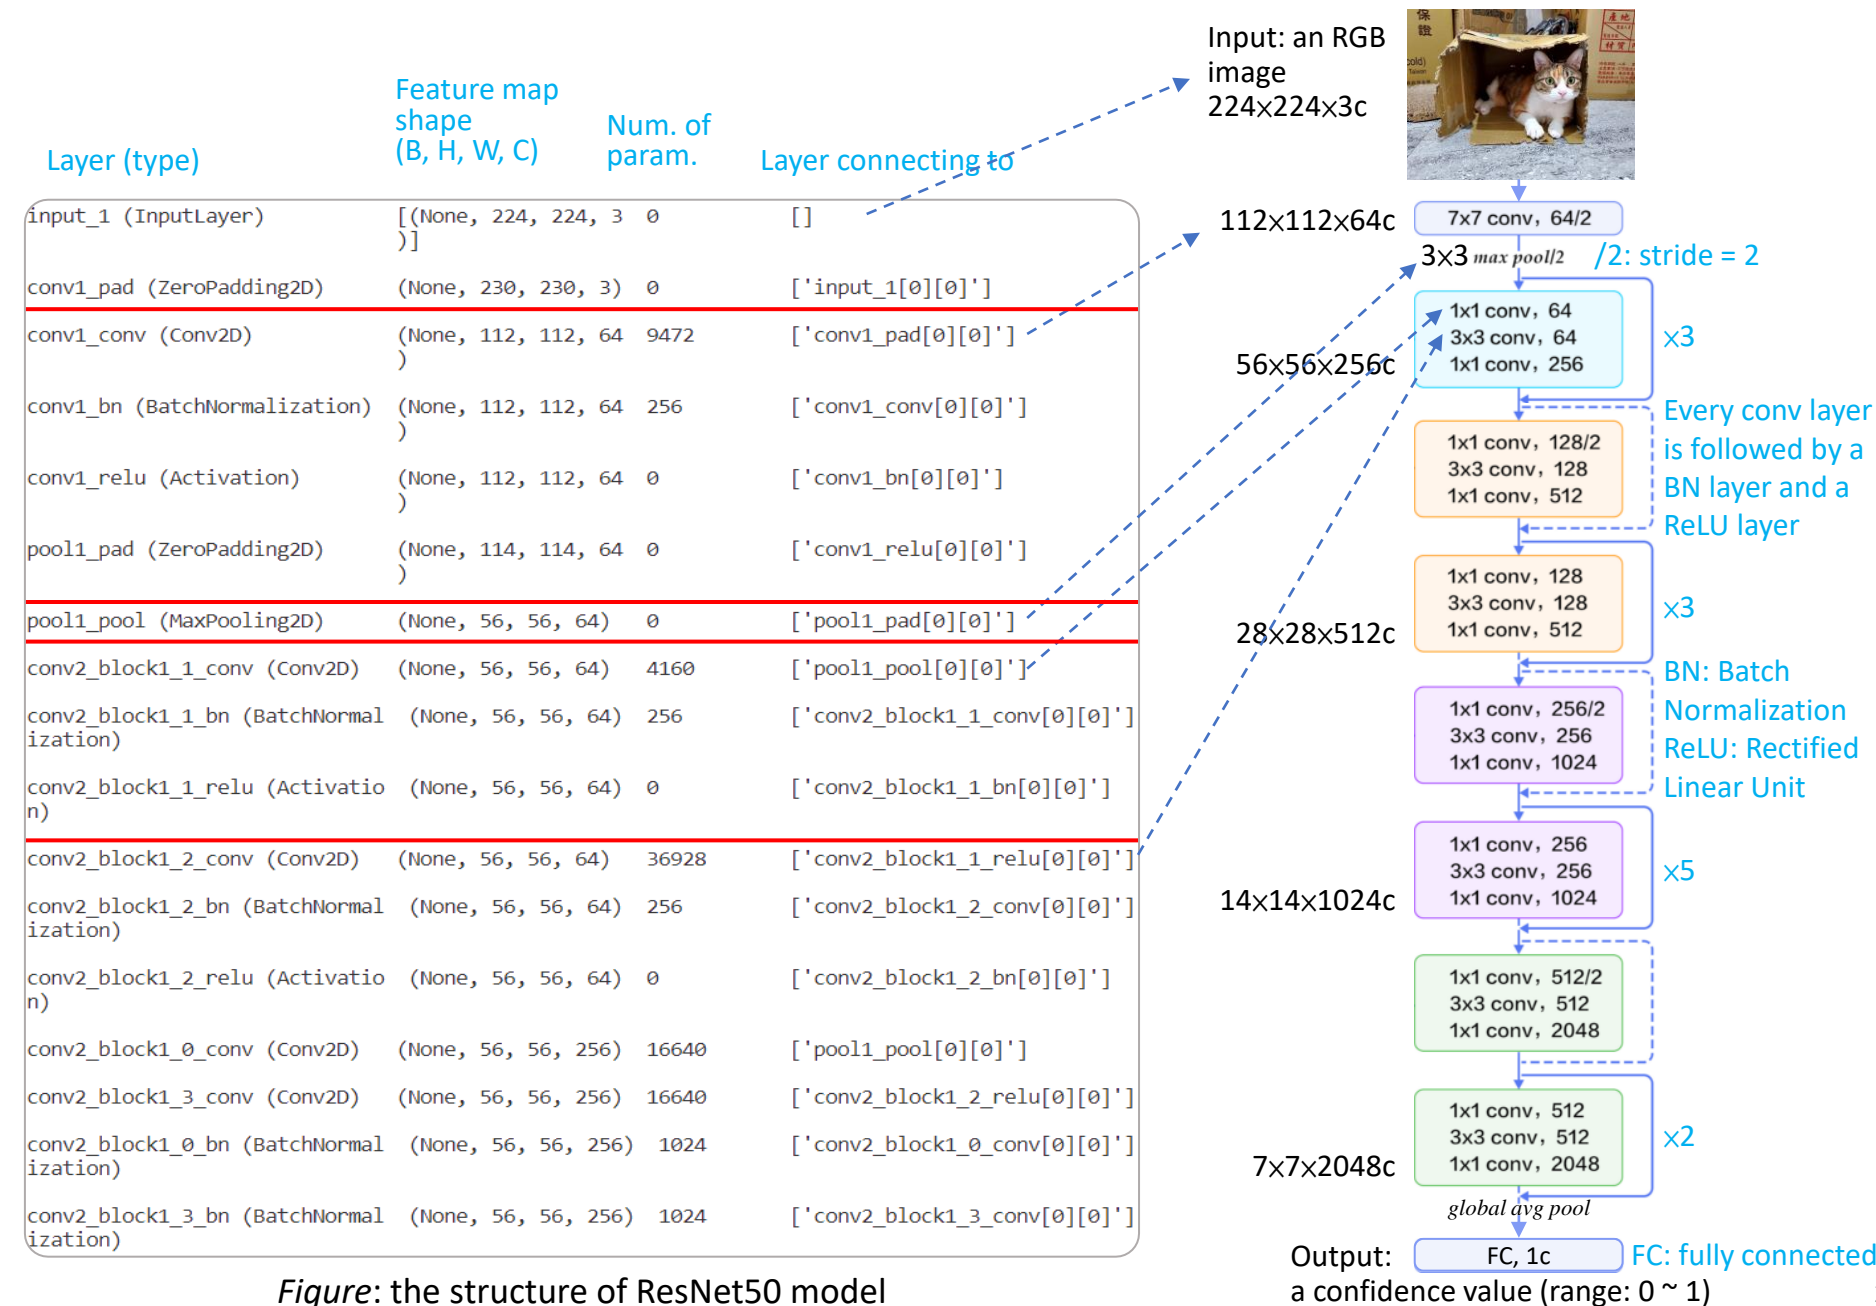
(1) Click the button **“3. Show Model Structure”**

(2) Run the function to show the structure **in the terminal**



5.2 (5%) Show the structure of ResNet50 model

(出題：Shan)



5.3 (5%) Improve ResNet50 with Random-Erasing

- 1) At home: Set up **Random-Erasing** in codes for model training (train.py)
 - (1) Train **2 ResNet50 models** with training dataset
 - Hint:(a) PyTorch ([tutorial](#)): write a for loop to validate the model
 - (a) With Random-Erasing
 - (b) Without Random-Erasing

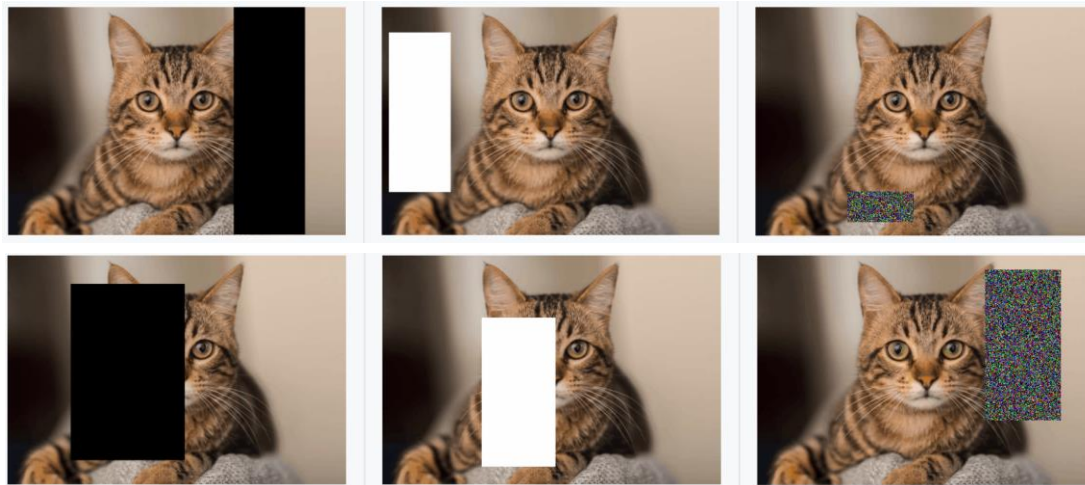


Figure1: Examples of the use of Random-Erasing

- 2) When the demo: Show your **codes** about Random-Erasing in train.py

```
transform = transforms.Compose([
    transforms.Resize(224),
    transforms.CenterCrop(224),
    transforms.RandomHorizontalFlip(),
    transforms.RandomVerticalFlip(),
    transforms.ToTensor(),
    transforms.RandomErasing(),
])
```

R. Reference

[Random Erasing Data Augmentation](#)

5.3 (5%) Compare the accuracies of 2 ResNet50 models on validation dataset

(出題：Shan)

1) At home:

(1) Validate 2 ResNet50 models with validation dataset

→ Hint:

(a) PyTorch ([tutorial](#)): write a for loop to validate the model

(2) Plot the accuracy values with a bar chart

(3) Save the figure

2) When the demo:

(1) Click the button “4. Show Comparison”

(2) Show the saved figure of accuracy comparison in a new window

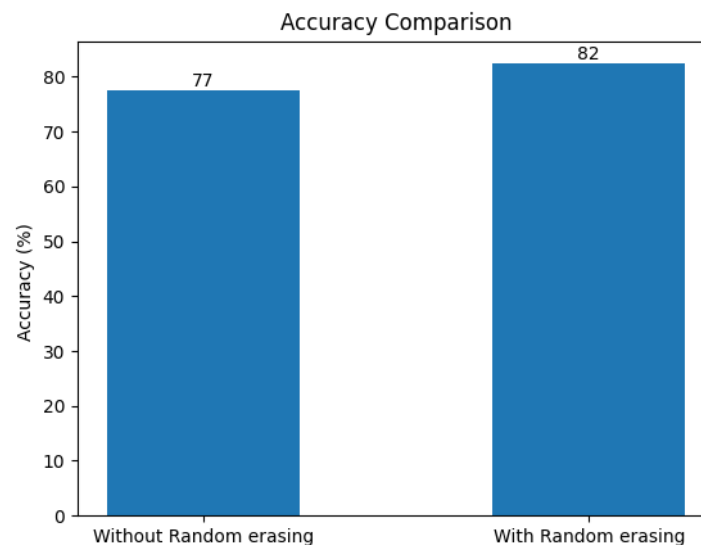
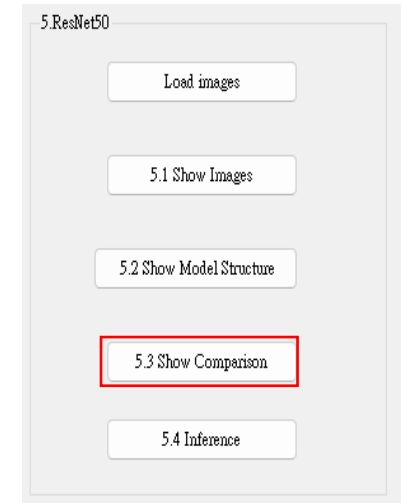


Figure1: Accuracy Comparison

Notice: this is an example, the numbers might differ

5.4 (5%) Use the better-trained model to run inference and show the predicted class label

1) At home:

(1) Load the trained model

→ Hint:

(a) PyTorch: `torch.nn.Module.load_state_dict()`

(2) Click the button “Load Image” to select 1 image arbitrarily

→ Hint: `PyQt5.QtWidgets.QFileDialog.getOpenFileName()`

(3) Show the loaded image in the GUI

(4) Resize the loaded image to **224×224×3c** (RGB)

(5) Click the button “5. Inference” to run inference on the resized image

→ Hint:

(a) PyTorch: pass an image when calling `torch.nn.Module` object to run inference, ex: `trained_model(img)`

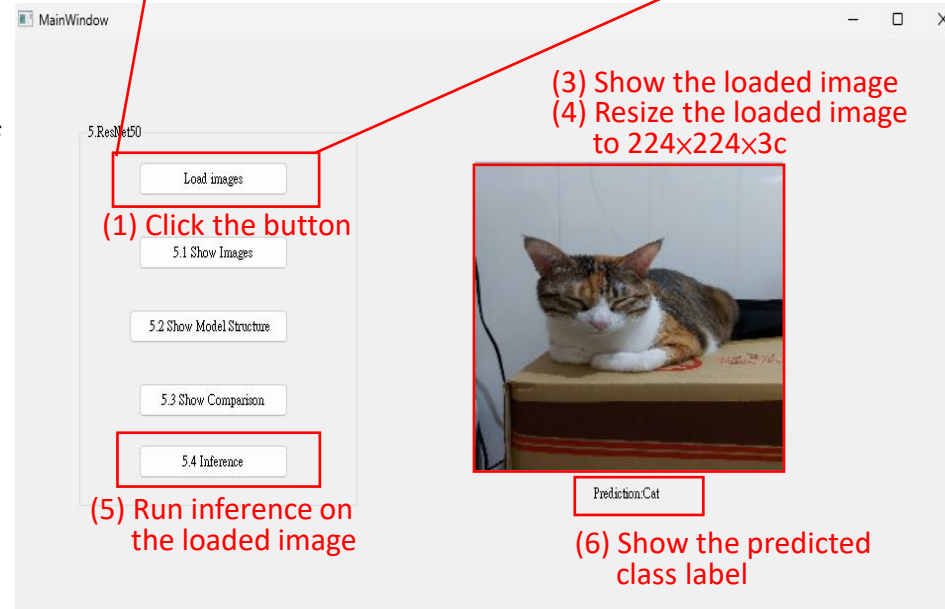
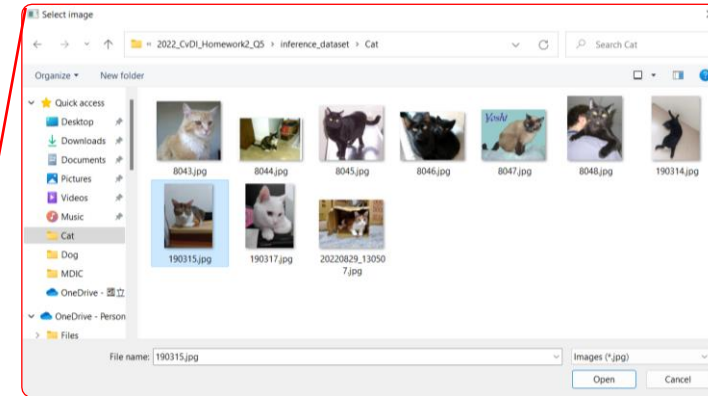
(6) Show the predicted class label

→ Hint: decide the class label with a **threshold** of the **output value**.

Ex: class label = $\begin{cases} \text{Cat, } output < thresh \\ \text{Dog, } output \geq thresh \end{cases}$
 $thresh = 0.5$

2) When the demo: **repeat** the process

(2) Select 1 image arbitrarily



5. Train a Cat-Dog Classifier Using ResNet50 – Demo Video

- This is an example illustrating the objectives from 5.1 ~ 5..

The screenshot displays a Python IDE (likely PyCharm) with a project named "GUI.py". The main.py file contains the following code:

```
12 self.file=File_UI()
13 self.model=Q5()
14 self.path=""
```

The IDE's console window shows the output of the training process, including the model architecture and parameters:

```
BatchNorm2d-149 [-1, 512, 7, 7] 1,024
ReLU-150 [-1, 512, 7, 7] 0
Conv2d-151 [-1, 2048, 7, 7] 1,048,576
BatchNorm2d-152 [-1, 2048, 7, 7] 4,096
ReLU-153 [-1, 2048, 7, 7] 0
ReLU-154 [-1, 2048, 7, 7] 0
Bottleneck-155 [-1, 2048, 7, 7] 0
AvgPool2d-156 [-1, 2048, 1, 1] 0
Linear-157 [-1, 1] 2,049
Sigmoid-158 [-1, 1] 0
=====
Total params: 17,650,241
Trainable params: 17,650,241
Non-trainable params: 0
=====
Input size (MB): 0.57
Forward/backward pass size (MB): 281.77
Params size (MB): 67.33
Estimated Total Size (MB): 349.67
=====
QCoreApplication::exec: The event loop is already running
Traceback (most recent call last):
File "C:\Users\User\anaconda3\envs\torch\lib\site-packages\PIL\fp.seek(0)
AttributeError: 'str' object has no attribute 'seek'
During handling of the above exception, another exception occurred
Traceback (most recent call last):
File "C:\Users\User\OneDrive - 國立成功大學 National Cheng Kung University\Homework\computer_vision_master\W2_5\main.py", line 42, in Test
result=self.model.Test(self.path)
File "C:\Users\User\OneDrive - 國立成功大學 National Cheng Kung University\Homework\computer_vision_master\W2_5\Q5.py", line 60, in Test
image=img.imread(path)
File "C:\Users\User\anaconda3\envs\torch\lib\site-packages\matplotlib\image.py", line 1541, in imread
with img_open(fname) as image:
File "C:\Users\User\anaconda3\envs\torch\lib\site-packages\PIL\image.py", line 3137, in open
fp = io.BytesIO(fp.read())
AttributeError: 'str' object has no attribute 'read'
PS C:\Users\User\OneDrive - 國立成功大學 National Cheng Kung University\Homework\computer_vision_master\W2_5>
c:; cd "c:\Users\User\OneDrive - 國立成功大學 National Cheng Kung University\Homework\computer_vision_master\W2_5"; & "C:\Users\User\anaconda3\envs\torch\python.exe" "c:\Users\User\vscode\extensions\ms-python.python-2023.20.0\pythonFiles\lib\python\debugpy\adapter\..\..\debugpy\launcher" "65378" "--" "C:\Users\User\OneDrive - 國立成功大學 National Cheng Kung University\Homework\computer_vision_master\W2_5\main.py"
```

The GUI window, titled "MainWindow", contains a "5.ResNet50" section with the following buttons:

- Load images
- 5.1 Show Images (highlighted)
- 5.2 Show Model Structure
- 5.3 Show Comparison
- 5.4 Inference

The "Prediction:" label is visible at the bottom right of the GUI window.